

Hello World, part 2

Agenda for today

Today we're going to go quickly back over loading data and then return to the topic of filtering, selecting and arranging data. We'll then turn to some calculations using the concepts of summarizing (self explanatory) and mutating (creating new variables).

Rmarkdown

As mentioned last time, an Rmarkdown file contains two basic elements: text and code. That text and code can be combined or “knitted” into a variety of different document formats. Lets get you started by creating your own Rmarkdown file and knitting it.

Load relevant libraries

```
library(tidyverse)
```

Load The Data

```
df<-readRDS("sc_debt.Rds")
names(df)
```

```
## [1] "unitid"      "instnm"      "stabbr"      "grad_debt_mdn"
## [5] "control"     "region"      "preddeg"     "openadmp"
## [9] "adm_rate"    "ccbasic"     "sat_avg"     "md_earn_wne_p6"
## [13] "ugds"        "costt4_a"    "selective"    "research_u"
```

Name	Definition
unitid	Unit ID
instnm	Institution Name
stabbr	State Abbreviation
grad_debt_mdn	Median Debt of Graduates
control	Control Public or Private
region	Census Region
preddeg	Predominant Degree Offered: Associates or Bachelors
openadmp	Open Admissions Policy: 1= Yes, 2=No,3=No 1st time students
adm_rate	Admissions Rate: proportion of applications accepted
ccbasic	Type of institution– see here
selective	Institution admits fewer than 10 % of applicants, 1=Yes, 0=No
research_u	Institution is a research university 1=Yes, 0=No
sat_avg	Average Sat Scores
md_earn_wne_p6	Average Earnings of Recent Graduates
ugds	Number of undergraduates

Looking at datasets

We can use “glimpse” to see what’s in a dataset. This gives a very quick rundown of the variables and the first few observations.

```
glimpse(df)

## Rows: 2,546
## Columns: 16
## $ unitid      <int> 100654, 100663, 100690, 100706, 100724, 100751, 100760, ~
## $ instnm     <chr> "Alabama A & M University", "University of Alabama at B~
## $ stabbr     <chr> "AL", "AL", "AL", "AL", "AL", "AL", "AL", "AL", "AL", "~
## $ grad_debt_mdn <int> 33375, 22500, 27334, 21607, 32000, 23250, 12500, 19500, ~
## $ control    <chr> "Public", "Public", "Private", "Public", "Public", "Pub~
## $ region     <chr> "Southeast", "Southeast", "Southeast", "Southeast", "So~
## $ preddeg    <chr> "Bachelor's", "Bachelor's", "Associate", "Bachelor's", ~
## $ openadmp   <int> 2, 2, 1, 2, 2, 1, NA, 2, 2, 2, 1, 1, 2, 1, 1, 2, 2, ~
## $ adm_rate   <dbl> 0.9175, 0.7366, NA, 0.8257, 0.9690, 0.8268, NA, NA, 0.9~
## $ ccbasic    <int> 18, 15, 20, 16, 19, 15, 2, 22, 18, 15, 21, 1, 5, 19, 7, ~
## $ sat_avg    <int> 939, 1234, NA, 1319, 946, 1261, NA, NA, 1082, 1300, 123~
## $ md_earn_wne_p6 <int> 25200, 35100, 30700, 36200, 22600, 37400, 23100, 33400, ~
## $ ugsd      <int> 5271, 13328, 365, 7785, 3750, 31900, 1201, 2677, 4407, ~
## $ costt4_a   <int> 23053, 24495, 14800, 23917, 21866, 29872, 10493, NA, 19~
## $ selective  <dbl> 0, 0, NA, 0, 0, 0, NA, NA, 0, 0, 0, NA, NA, 0, NA, NA, ~
## $ research_u <dbl> 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
```

Types of Variables

Notice that for each variable, it shows a different type, in angle brackets <>. So for instance, `instnm` has a type of <chr>. This is short for character— it’s also called a string variable.

Here are the types of data in this dataset

- <int> Integer data
- <chr> Character or string data
- <dbl> Double, (double-precision floating point) or just numeric data— can be measured down to an arbitrary number of data points.

This information is useful, because we wouldn’t want to try to run some kind of numeric analysis on string data. The average of institution names wouldn’t make a lot of sense (but it would probably be Southeast North State University College).

We’ll talk more about data types later, but we should also quickly note that there are some variables in this dataset where the numbers represent a characteristic, rather than a measurement. For instance, the variable `research_u` is set up—coded— such that a “1” indicates that the college is a research university and a “0” indicates that it is not a research university. The 1 and 0 don’t measure anything, they just indicate a characteristic.

Filter, Select, Arrange

Today, we’ll pick up where we left off— with the key commands of filter, select, and arrange.

In exploring data, many times we want to look at smaller parts of the dataset. There are three commands we’ll use today that help with this.

`-filter` selects only those cases or rows that meet some logical criteria.

-select selects only those variables or columns that meet some criteria

-arrange arranges the rows of a dataset in the way we want.

For more on these, please see this vignette.

We can look at the first 5 rows:

```
head(df)

## # A tibble: 6 x 16
##   unitid instnm   stabbr grad_debt_mdn control region preddeg openadmp adm_rate
##   <int> <chr>     <chr>         <int> <chr>   <chr> <chr>      <int>    <dbl>
## 1 100654 Alabama ~ AL             33375 Public  South~ Bachel~      2    0.918
## 2 100663 Universi~ AL             22500 Public  South~ Bachel~      2    0.737
## 3 100690 Amridge ~ AL             27334 Private South~ Associ~      1    NA
## 4 100706 Universi~ AL             21607 Public  South~ Bachel~      2    0.826
## 5 100724 Alabama ~ AL             32000 Public  South~ Bachel~      2    0.969
## 6 100751 The Univ~ AL             23250 Public  South~ Bachel~      2    0.827
## # ... with 7 more variables: ccbasic <int>, sat_avg <int>,
## #   md_earn_wne_p6 <int>, ugds <int>, costt4_a <int>, selective <dbl>,
## #   research_u <dbl>
```

Or the last 5 rows:

```
tail(df)

## # A tibble: 6 x 16
##   unitid instnm   stabbr grad_debt_mdn control region preddeg openadmp adm_rate
##   <int> <chr>     <chr>         <int> <chr>   <chr> <chr>      <int>    <dbl>
## 1 493716 Yeshiva ~ NJ              NA Private North~ Associ~      2    0.477
## 2 493725 Universi~ AR              NA Public  South~ Bachel~      1    NA
## 3 493822 College ~ RI              NA Private New E~ Bachel~      1    NA
## 4 494630 Christ M~ TX              NA Private South~ Bachel~      1    NA
## 5 494685 Urshan C~ MO              NA Private Plains Bachel~      2    0.836
## 6 494737 Yeshiva ~ NY              NA Private North~ Bachel~      1    NA
## # ... with 7 more variables: ccbasic <int>, sat_avg <int>,
## #   md_earn_wne_p6 <int>, ugds <int>, costt4_a <int>, selective <dbl>,
## #   research_u <dbl>
```

Using filter in combination with other commands

filter can be used with any command that retruns true or false. This can be really powerful, for instance the command `str_detect` “detects” the relevant string in the data, so we can look for any college with the word “Colorado” in its name.

```
df%>%
  filter(str_detect(instnm,"Colorado"))%>%
  select(instnm,adm_rate,sat_avg)
```

```
## # A tibble: 12 x 3
##   instnm                                adm_rate sat_avg
##   <chr>                                <dbl>    <int>
## 1 University of Colorado Denver/Anschutz Medical Campus  0.673    1124
## 2 University of Colorado Colorado Springs                0.872    1136
## 3 University of Colorado Boulder                        0.784    1276
## 4 Colorado Christian University                        NA        NA
```

```
## 5 Colorado College 0.135 NA
## 6 Colorado School of Mines 0.531 1342
## 7 Colorado State University-Fort Collins 0.814 1204
## 8 Colorado Mesa University 0.782 1063
## 9 University of Northern Colorado 0.908 1096
## 10 Colorado State University Pueblo 0.930 1047
## 11 Western Colorado University 0.842 1114
## 12 Colorado State University-Global Campus 0.986 1048
```

We can combine this with the `|` operator, which remember stands for “or.” Let’s say we want all the institutions in Colorado OR California.

```
df%>%
  filter(str_detect(instnm,"Colorado") | str_detect(instnm,"California"))%>%
  select(instnm,adm_rate,sat_avg)
```

```
## # A tibble: 57 x 3
##   instnm adm_rate sat_avg
##   <chr>   <dbl>   <int>
## 1 California Institute of Integral Studies NA NA
## 2 California Baptist University 0.783 1096
## 3 California College of the Arts 0.850 NA
## 4 California Institute of Technology 0.0642 1557
## 5 California Lutheran University 0.714 1168
## 6 California Polytechnic State University-San Luis Obispo 0.284 1342
## 7 California State University-Bakersfield 0.807 NA
## 8 California State University-Stanislaus 0.893 NA
## 9 California State University-San Bernardino 0.686 985
## 10 California State Polytechnic University-Pomona 0.546 1143
## # ... with 47 more rows
```

We can also put this together in one (notice that everything goes inside the quotes)

```
df%>%
  filter(str_detect(instnm,"Colorado|California"))%>%
  select(instnm,adm_rate,sat_avg)
```

```
## # A tibble: 57 x 3
##   instnm adm_rate sat_avg
##   <chr>   <dbl>   <int>
## 1 California Institute of Integral Studies NA NA
## 2 California Baptist University 0.783 1096
## 3 California College of the Arts 0.850 NA
## 4 California Institute of Technology 0.0642 1557
## 5 California Lutheran University 0.714 1168
## 6 California Polytechnic State University-San Luis Obispo 0.284 1342
## 7 California State University-Bakersfield 0.807 NA
## 8 California State University-Stanislaus 0.893 NA
## 9 California State University-San Bernardino 0.686 985
## 10 California State Polytechnic University-Pomona 0.546 1143
## # ... with 47 more rows
```

Reminder: logical operators

Here are (many of) the logical operators that we use in R:

- >, <: greater than, less than
- >=, <=: greater than or equal to, less than or equal to
- ! :not, as in != not equal to
- & AND
- | OR

Quick Exercise Select colleges that are from Texas AND have the word “community” in their name (the name variable is `instnm`).

INSERT CODE HERE

Extending Select

Select can also be used with other characteristics.

For quick guide on this: <https://dplyr.tidyverse.org/reference/select.html>

For example, we can select just variables that contain the word “region”

```
df%>%
  select(contains("region"))
```

```
## # A tibble: 2,546 x 1
##   region
##   <chr>
## 1 Southeast
## 2 Southeast
## 3 Southeast
## 4 Southeast
## 5 Southeast
## 6 Southeast
## 7 Southeast
## 8 Southeast
## 9 Southeast
## 10 Southeast
## # ... with 2,536 more rows
```

We can also select just variables that

```
df%>%
  select(where(is.numeric))
```

```
## # A tibble: 2,546 x 11
##   unitid grad_debt_mdn openadmp adm_rate ccbasic sat_avg md_earn_wne_p6 ugds
##   <int>   <int>      <int>   <dbl>   <int>   <int>      <int> <int>
## 1 100654   33375        2    0.918     18     939     25200  5271
## 2 100663   22500        2    0.737     15    1234     35100 13328
## 3 100690   27334        1    NA        20     NA     30700   365
## 4 100706   21607        2    0.826     16    1319     36200  7785
## 5 100724   32000        2    0.969     19     946     22600  3750
## 6 100751   23250        2    0.827     15    1261     37400 31900
## 7 100760   12500        1    NA         2     NA     23100  1201
## 8 100812   19500       NA    NA        22     NA     33400  2677
## 9 100830   24826        2    0.904     18    1082     30100  4407
## 10 100858   21281        2    0.807     15    1300     39500 24209
## # ... with 2,536 more rows, and 3 more variables: costt4_a <int>,
## #   selective <dbl>, research_u <dbl>
```

Quick Exercise Use the same setup to select only character variables (`is.character`)

```
# INSERT CODE HERE
```

Summarizing Data

To summarize data, we use the `summarize` command. Inside that command, we tell R two things: what to call the new object (a data frame, really) that we're creating, and what numerical summary we would like. The code below summarizes median debt for the colleges in the dataset by calculating the average of median debt for all institutions.

Notice that inside the `mean` command

```
df%>%
  summarize(mean_debt=mean(grad_debt_mdn,na.rm=TRUE))
```

```
## # A tibble: 1 x 1
##   mean_debt
##       <dbl>
## 1    19646.
```

Quick Exercise Summarize the average entering SAT scores in this dataset.

```
# INSERT CODE HERE
```

Combining Commands

We can also combine commands, so that summaries are done on only a part of the dataset. Below, we summarize median debt for selective schools, and not very selective schools.

```
df%>%
  filter(stabbr=="CA")%>%
  summarize(mean_adm_rate=mean(adm_rate,na.rm=TRUE))
```

```
## # A tibble: 1 x 1
##   mean_adm_rate
##       <dbl>
## 1         0.592
```

Quick Exercise Calculate average earnings for schools where SAT>1200 & the admissions rate is between 10 and 20 percent.

```
# INSERT CODE HERE
```

Mutate

`mutate` is the verb for changing variables in R. Let's say we want to create a variable that's set to 1 if the college admits less than 10 percent of the students who apply.

```
df<-df%>%
  mutate(selective=ifelse(adm_rate<=.1,1,0))
```

Or what if we want to create another new variable that changes the admissions rate from its current proportion to a percent?

```
df<-df%>%
  mutate(adm_rate_pct=adm_rate*100)
```

To figure out if that worked we can use `summarize`

```
df%>%
  summarize(mean_adm_rate_pct=mean(adm_rate_pct,na.rm=TRUE))

## # A tibble: 1 x 1
##   mean_adm_rate_pct
##               <dbl>
## 1                67.9
```

Quick Exercise Create a new variable that's set to 1 if the college has more than 10,000 undergraduate students

```
# INSERT CODE HERE
```


Plotting Data

Plotting data is another key skill as we're getting started. We're going to do a LOT more plotting, but let's just quickly look at the relationship between two variables using a plot. The basic tool for plotting in r is `ggplot`. To create a plot, we specify what goes on the x and y axis in an "aesthetics" argument, which is called `aes`. We then say what we want to draw on the plot– in this case a point for each case, which means we want to use `geom_point`.

Below is a plot of earnings of by SAT scores

```
## Plotting: bivariate
df%>%
  ggplot(aes(x=sat_avg,y=md_earn_wne_p6))+
  geom_point()
```

```
## Warning: Removed 1348 rows containing missing values (geom_point).
```



Quick Exercise Plot earnings by admission rate only for schools in California.

```
# INSERT CODE HERE
```

Quick Exercise Replicate the above plots, but put debt level on the y axis.

```
# INSERT CODE HERE
```